

# Final Project

2025-03-02

## 1. Data Preparing

```
# Load required libraries  
library(quantmod)
```

```
## Loading required package: xts
```

```
## Loading required package: zoo
```

```
##
```

```
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      as.Date, as.Date.numeric
```

```
## Loading required package: TTR
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
##      method          from
```

```
##      as.zoo.data.frame zoo
```

```
library(forecast)  
library(ggplot2)  
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats    1.0.0      v stringr    1.5.1
```

```
## v lubridate  1.9.3      v tibble     3.2.1
```

```
## v purrr      1.0.2      v tidyr      1.3.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::first()  masks xts::first()
```

```
## x dplyr::lag()    masks stats::lag()
```

```
## x dplyr::last()   masks xts::last()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(fracdiff)
library(vars)
```

```
## Loading required package: MASS
##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##   select
##
## Loading required package: strucchange
## Loading required package: sandwich
##
## Attaching package: 'strucchange'
##
## The following object is masked from 'package:stringr':
##
##   boundary
##
## Loading required package: urca
## Loading required package: lmtest
```

```
library(Metrics)
```

```
##
## Attaching package: 'Metrics'
##
## The following object is masked from 'package:forecast':
##
##   accuracy
```

```
library(dplyr)
library(lmtest)
library(tseries)
library(FinTS)
```

```
##
## Attaching package: 'FinTS'
##
## The following object is masked from 'package:forecast':
##
##   Acf
```

```
library(rugarch)
```

```
## Loading required package: parallel
##
## Attaching package: 'rugarch'
##
## The following object is masked from 'package:purrr':
```

```
##
##      reduce
##
## The following object is masked from 'package:stats':
##
##      sigma
```

```
cpi_data <- read.csv("CPIAUCNS.csv")
cpi_data$DATE <- as.Date(cpi_data$observation_date)
cpi_ts <- ts(cpi_data$CPIAUCNS, start = c(year(min(cpi_data$DATE)), month(min(cpi_data$DATE))), frequency = 12)

# Download data
getSymbols(c("DFF", "FEDFUNDS", "UNRATE", "GDP", "M2SL"), src = "FRED")
```

```
## [1] "DFF"      "FEDFUNDS" "UNRATE"    "GDP"      "M2SL"
```

```
# Convert to time series format, extracting data after 2000-01-01
fedfunds_ts <- window(ts(FEDFUNDS, start = c(1954, 7), frequency = 12), start = c(1999, 12), end = c(2025, 1))
unrate_ts <- window(ts(UNRATE, start = c(1948, 1), frequency = 12), start = c(1999, 12), end = c(2025, 1))
# gdp_ts <- window(ts(GDP, start = c(1947, 1), frequency = 4), start = c(2000, 1)) # Quarterly
m2_ts <- window(ts(M2SL, start = c(1959, 1), frequency = 12), start = c(1999, 12), end = c(2025, 1))

# Merge all time series into a dataset
dataset <- data.frame(
  Date = seq(as.Date("2000-01-01"), by = "month", length.out = length(cpi_ts)),
  CPI = as.numeric(cpi_ts),
  FEDFUNDS = as.numeric(fedfunds_ts[-1]),
  UNRATE = as.numeric(unrate_ts[-1]),
  M2 = as.numeric(m2_ts[-1])
)

# Create lagged versions of external regressors (1-month lag)
dataset_lag <- data.frame(
  Date = seq(as.Date("2000-01-01"), by = "month", length.out = length(cpi_ts)),
  CPI = as.numeric(cpi_ts),
  FEDFUNDS = as.numeric(fedfunds_ts[-length(fedfunds_ts)]),
  UNRATE = as.numeric(unrate_ts[-length(unrate_ts)]),
  M2 = as.numeric(m2_ts[-length(m2_ts)])
)
```

```
summary(dataset)
```

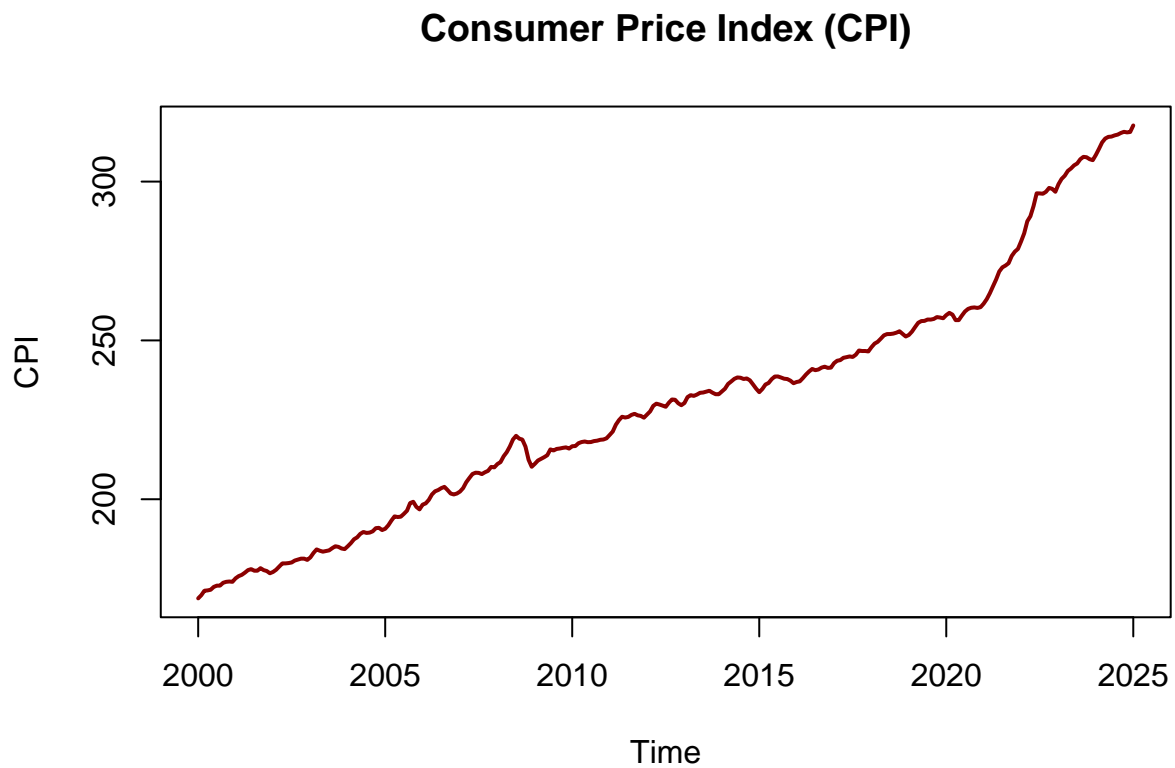
```
##      Date      CPI      FEDFUNDS      UNRATE
## Min.   :2000-01-01  Min.   :168.8  Min.   :0.050  Min.   : 3.40
## 1st Qu.:2006-04-01  1st Qu.:201.5  1st Qu.:0.140  1st Qu.: 4.20
## Median :2012-07-01  Median :229.8  Median :1.220  Median : 5.10
## Mean   :2012-07-01  Mean   :230.2  Mean   :1.928  Mean   : 5.69
## 3rd Qu.:2018-10-01  3rd Qu.:252.0  3rd Qu.:3.620  3rd Qu.: 6.40
## Max.   :2025-01-01  Max.   :317.7  Max.   :6.540  Max.   :14.80
##      M2
## Min.   : 4668
## 1st Qu.: 6807
```

```
## Median :10066
## Mean   :11399
## 3rd Qu.:14241
## Max.   :21750
```

## 2. EDA

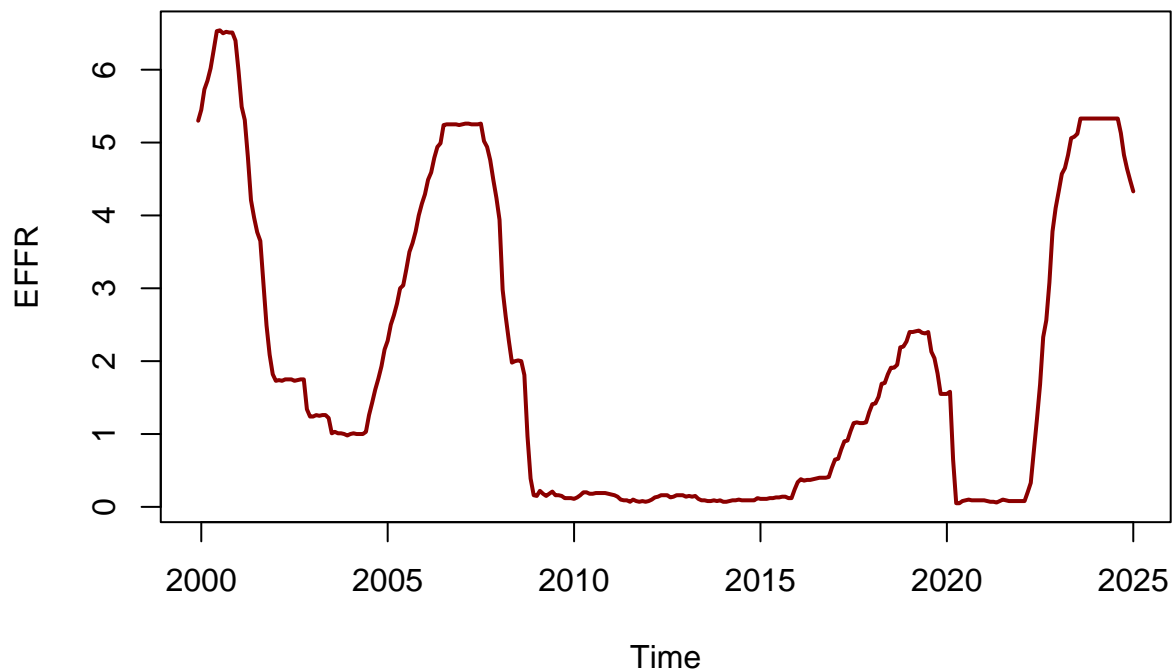
### (1) Time Series Plots

```
# Plot time series trends
plot(cpi_ts, main = "Consumer Price Index (CPI)",
     ylab = "CPI", xlab = "Time", col = "darkred", lwd = 2)
```



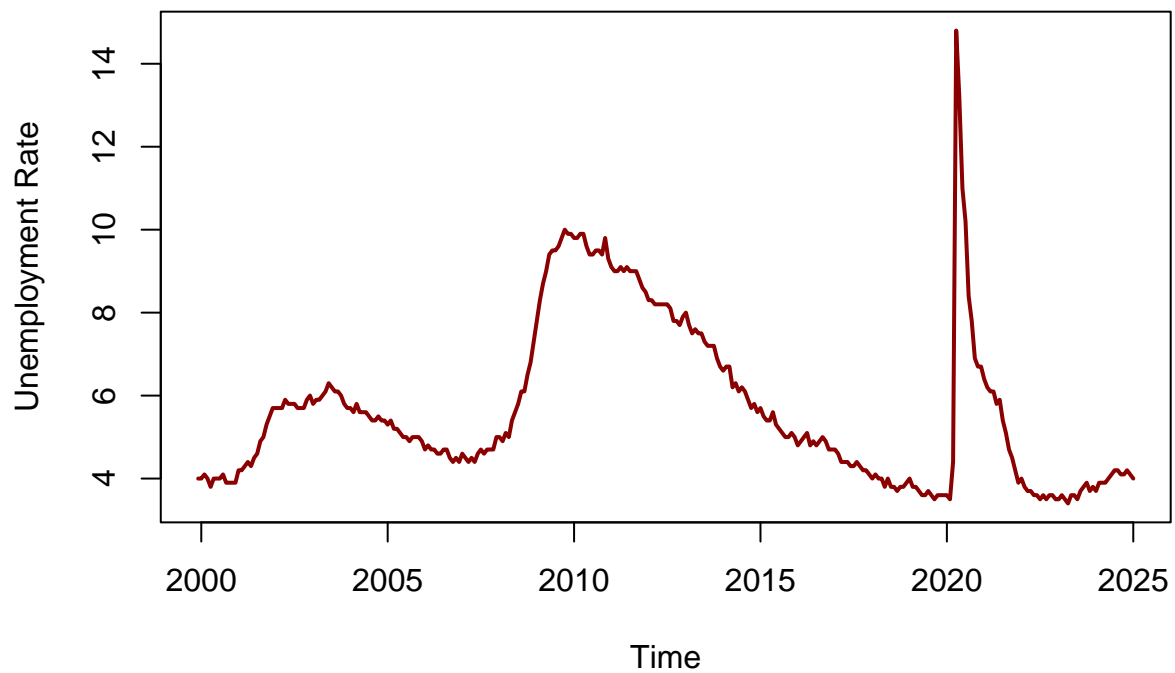
```
plot(fedfunds_ts, main = "Effective Federal Funds Rate",
     ylab = "EFFR", xlab = "Time", col = "darkred", lwd = 2)
```

## Effective Federal Funds Rate

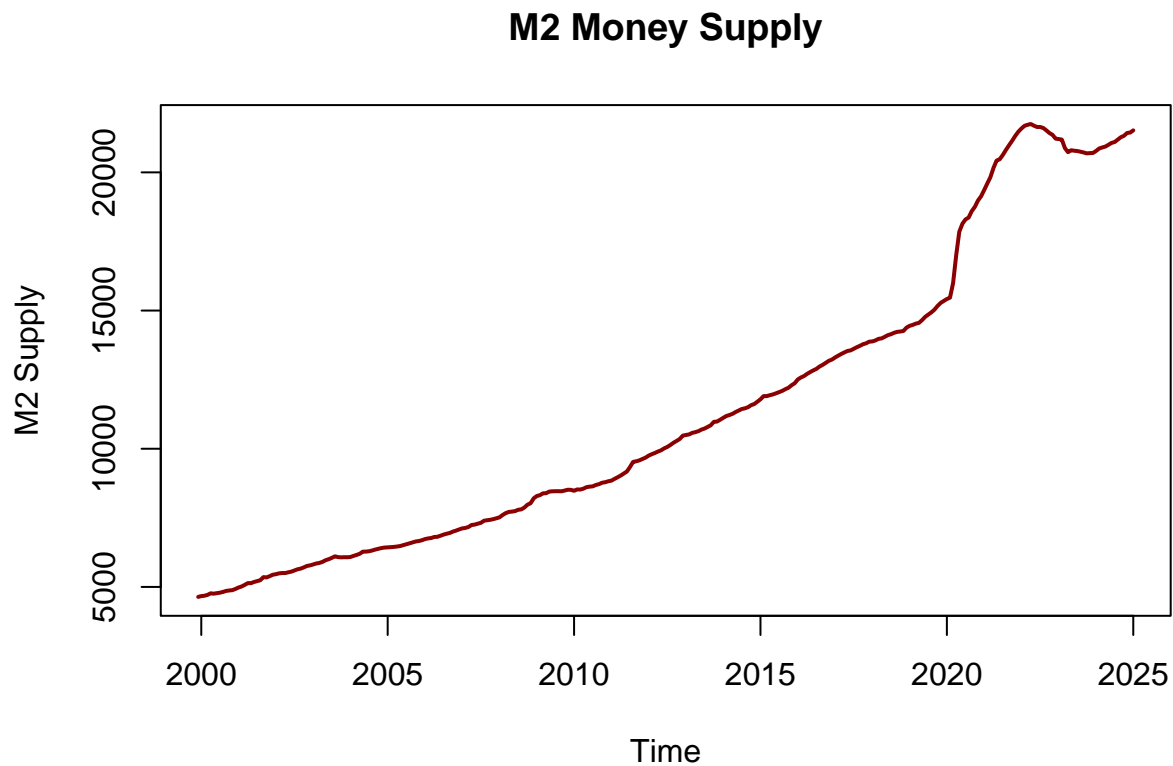


```
plot(unrate_ts, main = "Unemployment Rate",  
     ylab = "Unemployment Rate", xlab = "Time", col = "darkred", lwd = 2)
```

## Unemployment Rate

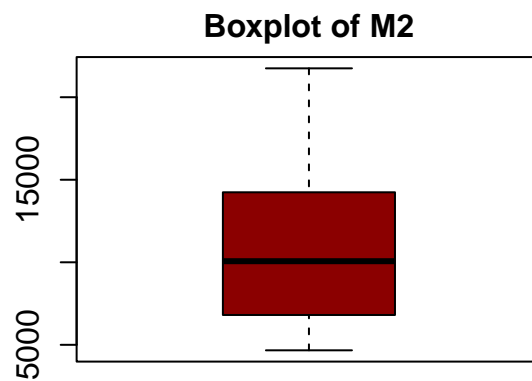
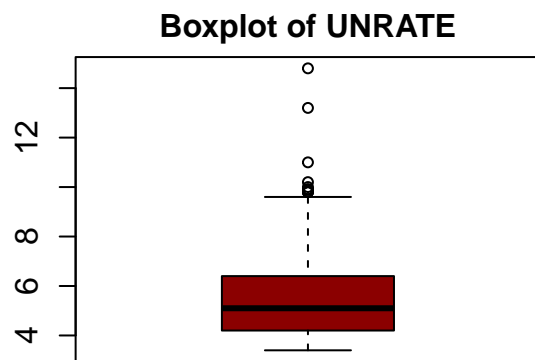
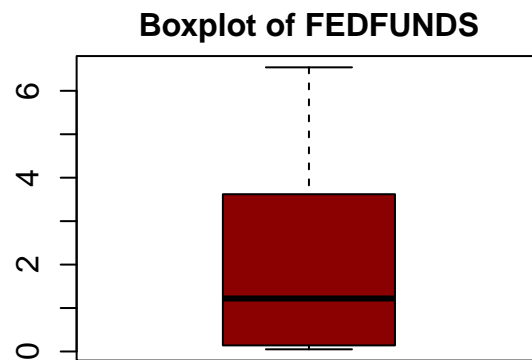
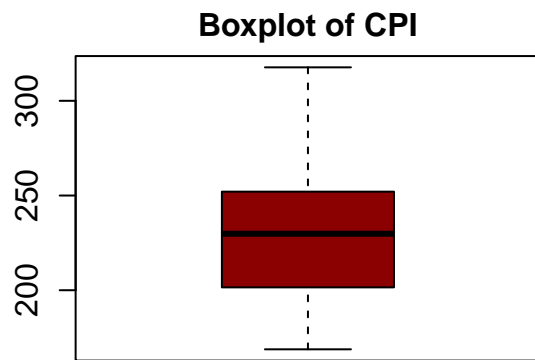


```
plot(m2_ts, main = "M2 Money Supply",
     ylab = "M2 Supply", xlab = "Time", col = "darkred", lwd = 2)
```



## (2) Boxplots

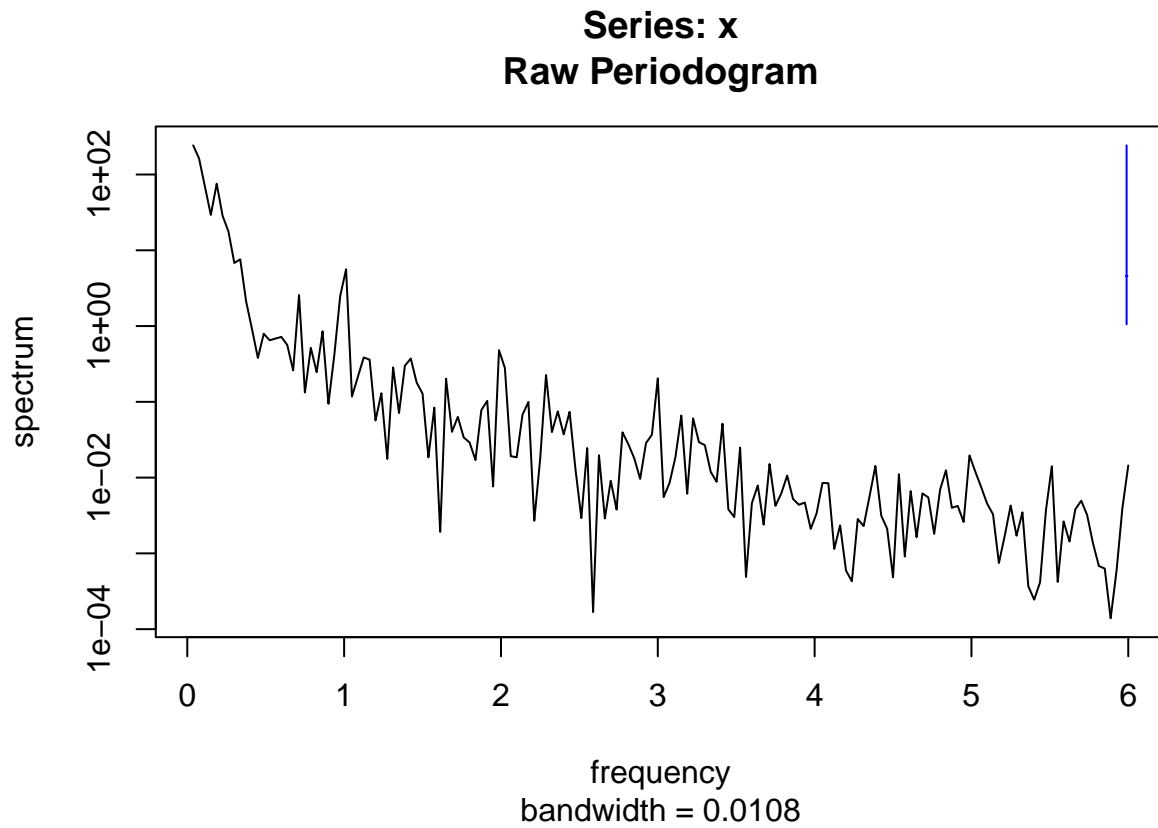
```
# Boxplots for each feature
par(mfrow = c(2,2), mar = c(2,3,2,2), cex.lab = 1.2, cex.axis = 1.2, cex.main = 1.2) # Arrange plots in a 2x2 grid
boxplot(dataset$CPI, main = "Boxplot of CPI", col = "darkred")
boxplot(dataset$FEDFUNDS, main = "Boxplot of FEDFUNDS", col = "darkred")
boxplot(dataset$UNRATE, main = "Boxplot of UNRATE", col = "darkred")
boxplot(dataset$M2, main = "Boxplot of M2", col = "darkred")
```



```
par(mfrow = c(1,1)) # Reset layout
```

### (3) Spectrum Analysis

```
# spectrum analysis  
spectrum(cpi_ts)
```

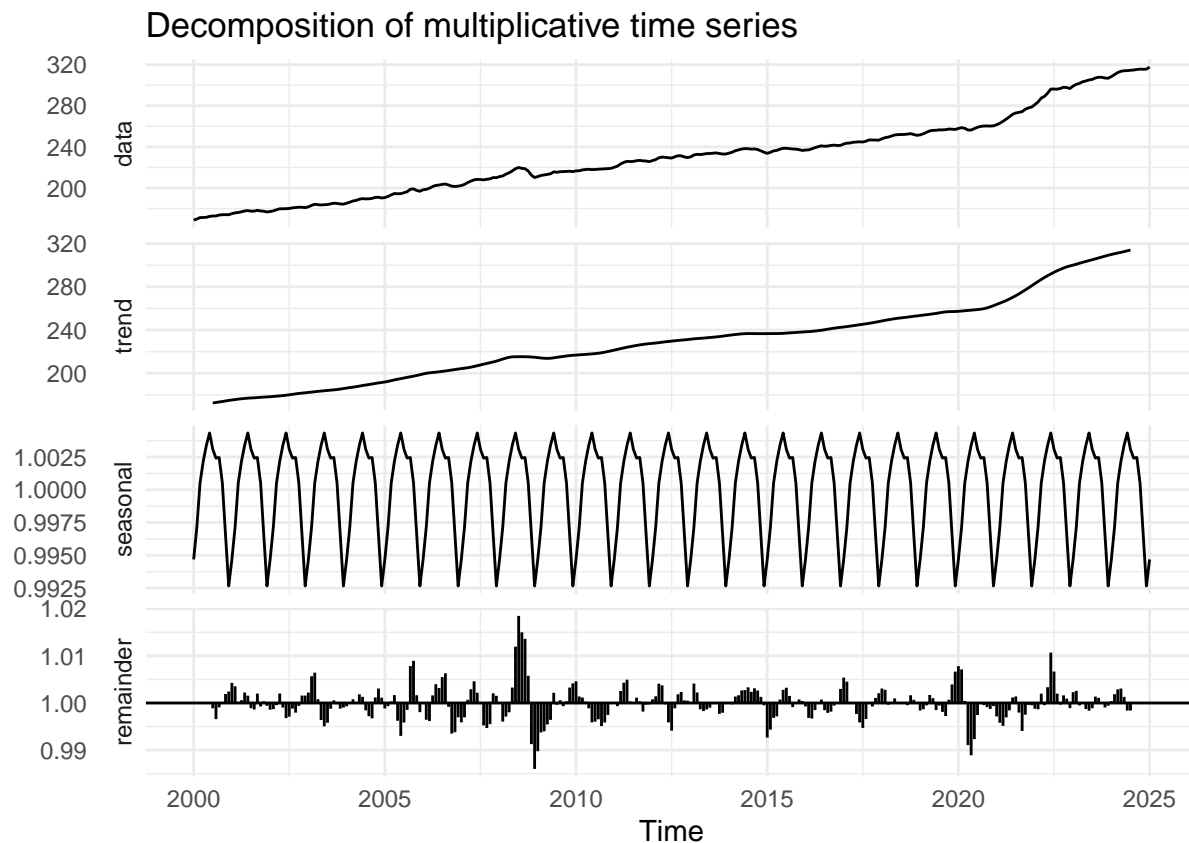


#### (4) Decomposition

```
# Decompose the time series into trend, seasonal, and random components
decomp <- decompose(cpi_ts, type = "multiplicative")

# Plot the decomposition
autoplot(decomp)+theme_minimal()
```





## (5) Check Stationarity

```
library(tseries)
# ADF Test for stationarity
adf_test <- adf.test(cpi_ts)
```

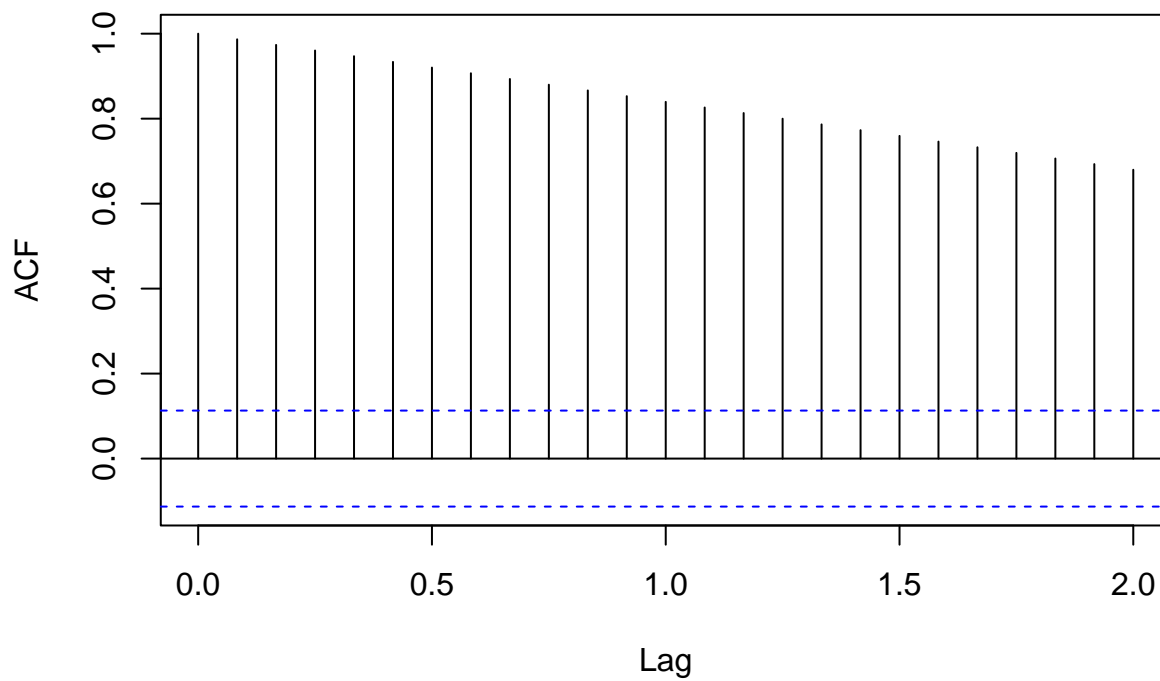
```
## Warning in adf.test(cpi_ts): p-value greater than printed p-value
```

```
print(adf_test)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: cpi_ts
## Dickey-Fuller = 0.59021, Lag order = 6, p-value = 0.99
## alternative hypothesis: stationary
```

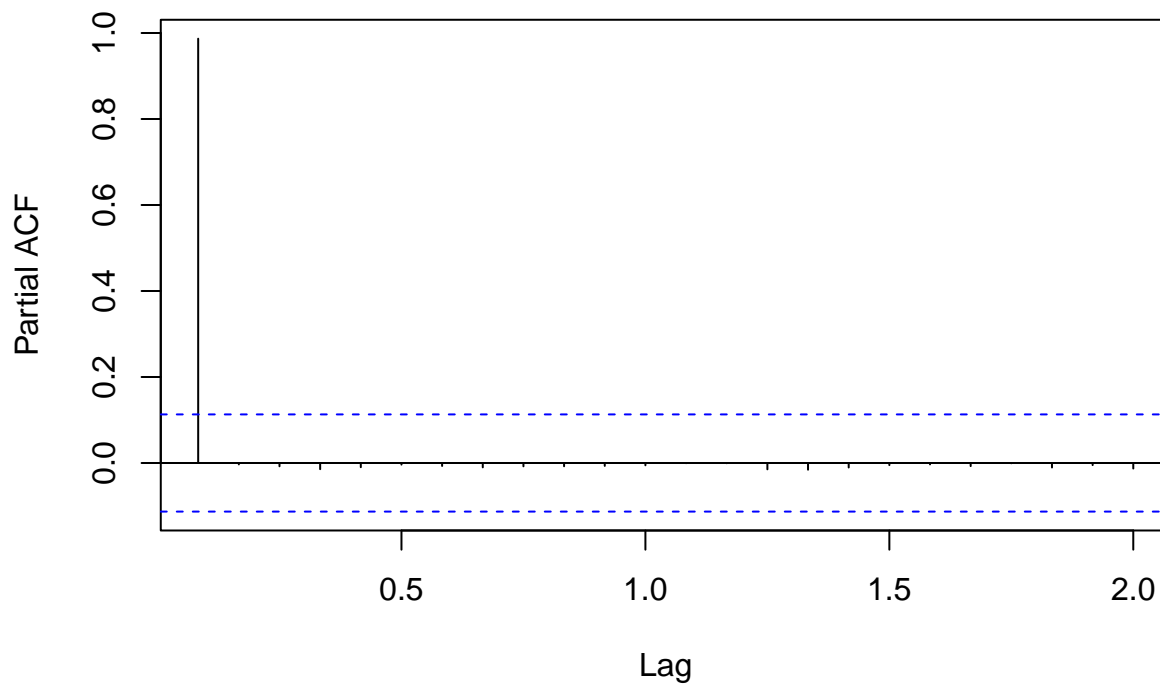
```
# Plot ACF and PACF
acf(cpi_ts, main = "ACF")
```

## ACF



```
pacf(cpi_ts, main = "PACF")
```

## PACF



The data is non-stationary because P-value = 0.99 that significantly greater than 0.05.

### 3. Models

#### (1) Univariate

```

cpi_ts_train <- window(cpi_ts, start = c(2000, 1), end = c(2023, 12))
cpi_ts_test  <- window(cpi_ts, start = c(2024, 1))

evaluate_performance <- function(actual, predicted) {
  mse_value <- mse(actual, predicted)
  rmse_value <- rmse(actual, predicted)
  mae_value <- mae(actual, predicted)
  rss <- sum((actual - predicted)^2) # Residual Sum of Squares
  tss <- sum((actual - mean(actual))^2) # Total Sum of Squares
  r_squared <- 1 - (rss/tss)
  results <- list(
    MSE = mse_value,
    RMSE = rmse_value,
    MAE = mae_value,
    R_squared = r_squared
  )
  return(results)
}
```

##### a. Regression Model

```

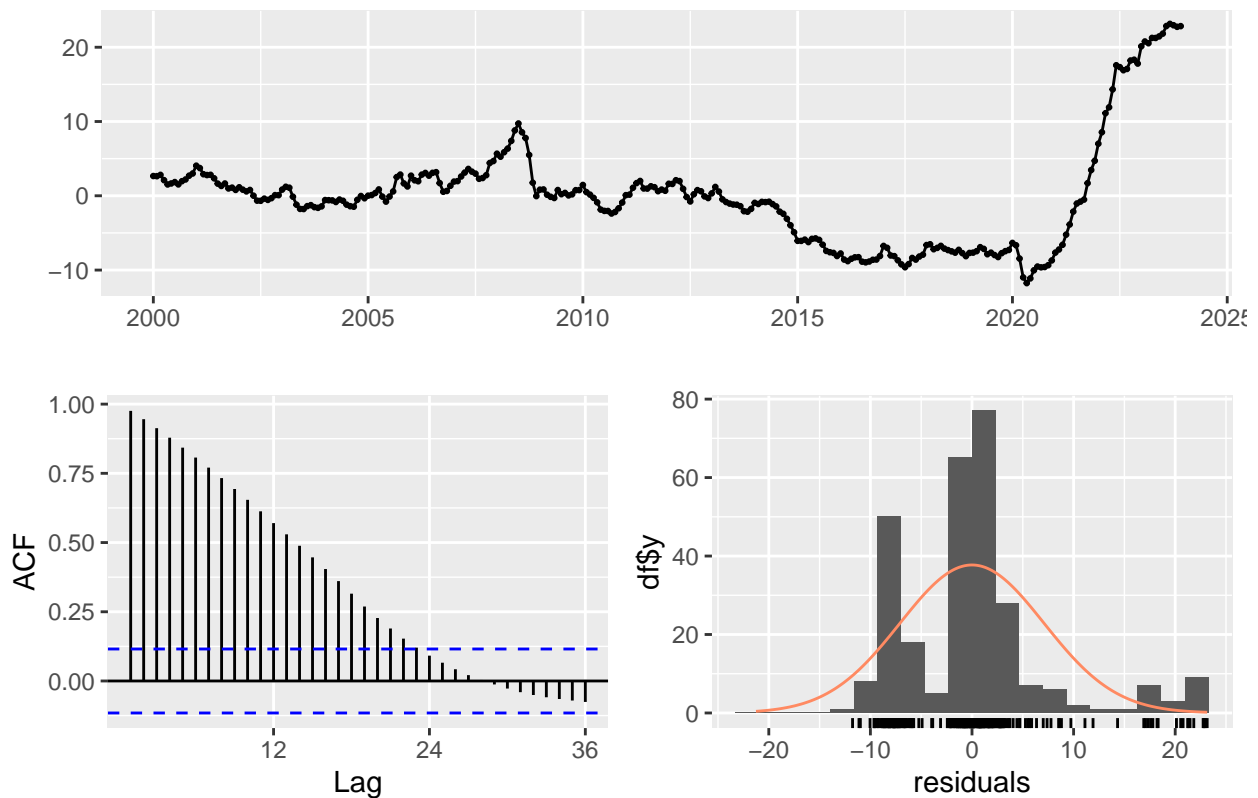
regression <- tslm(cpi_ts_train ~ trend + season)
summary(regression)
```

```
##
## Call:
## tslm(formula = cpi_ts_train ~ trend + season)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11.7568  -5.9115  -0.0911   1.9671  23.1474
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.657e+02  1.641e+00  100.968  <2e-16 ***
## trend        4.091e-01  5.136e-03   79.642  <2e-16 ***
## season2      6.107e-01  2.090e+00   0.292    0.770
## season3      1.430e+00  2.090e+00   0.684    0.494
## season4      1.832e+00  2.090e+00   0.876    0.382
## season5      2.196e+00  2.090e+00   1.051    0.294
## season6      2.555e+00  2.090e+00   1.222    0.223
## season7      2.362e+00  2.090e+00   1.130    0.260
## season8      2.264e+00  2.091e+00   1.083    0.280
## season9      2.324e+00  2.091e+00   1.111    0.267
## season10     1.981e+00  2.091e+00   0.948    0.344
## season11     1.165e+00  2.091e+00   0.557    0.578
```

```
## season12    3.589e-01  2.091e+00  0.172    0.864
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.241 on 275 degrees of freedom
## Multiple R-squared:  0.9586, Adjusted R-squared:  0.9567
## F-statistic:   530 on 12 and 275 DF,  p-value: < 2.2e-16
```

```
checkresiduals(regression)
```

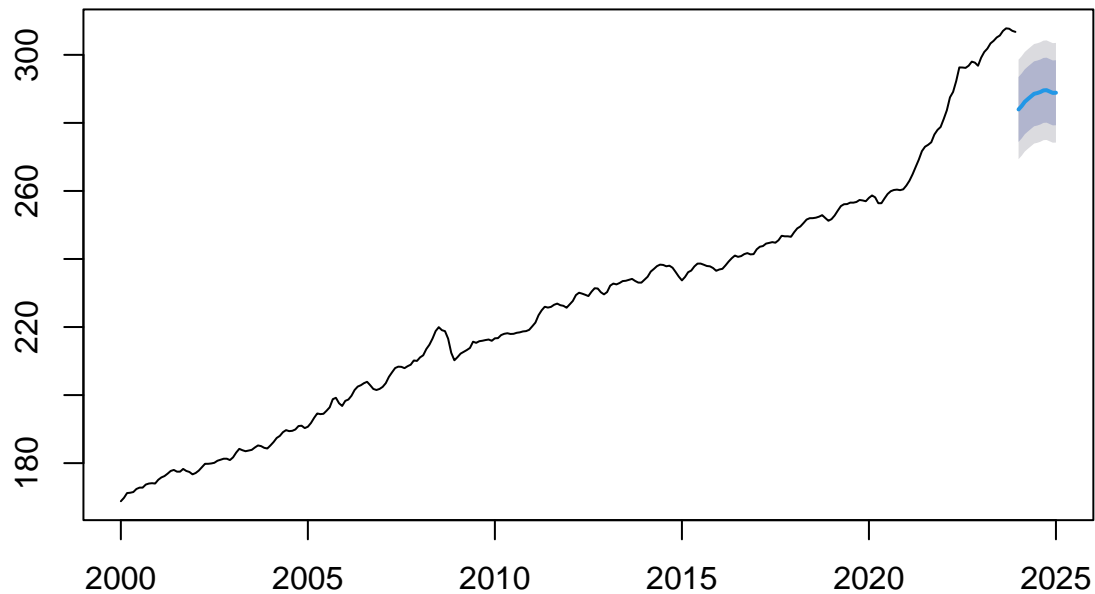
### Residuals from Linear regression model



```
##
## Breusch-Godfrey test for serial correlation of order up to 24
##
## data: Residuals from Linear regression model
## LM test = 285.97, df = 24, p-value < 2.2e-16
```

```
regression_forecast <- forecast(regression, h = 13)
plot(regression_forecast)
```

## Forecasts from Linear regression model



```
evaluate_performance(cpi_ts_test, regression_forecast$mean)
```

```
## $MSE
## [1] 683.1153
##
## $RMSE
## [1] 26.13648
##
## $MAE
## [1] 26.11893
##
## $R_squared
## [1] -123.4503
```

### b. ARIMA Model

```
lambda <- BoxCox.lambda(cpi_ts_train)
print(lambda)
```

```
## [1] -0.9999242
```

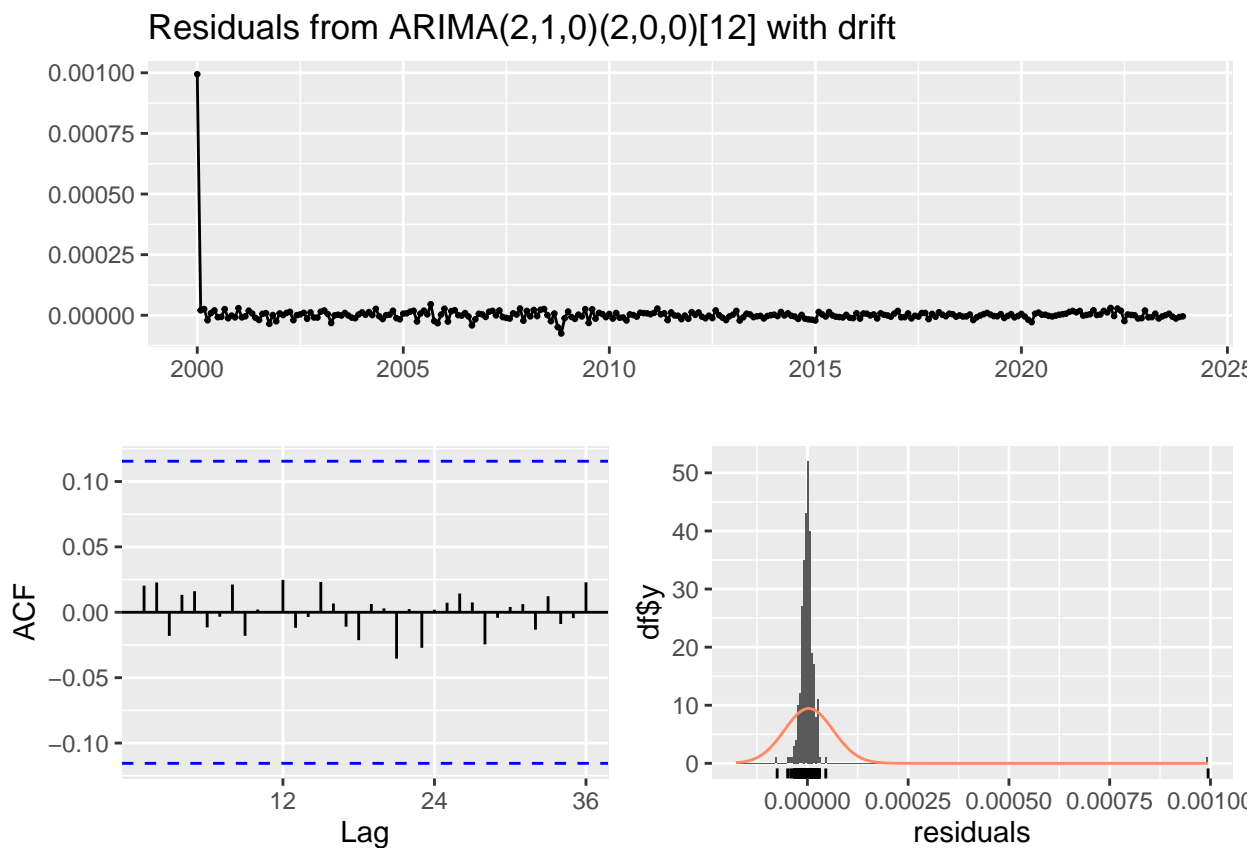
```
shapiro.test(cpi_ts_train)
```

```
##
## Shapiro-Wilk normality test
##
## data:  cpi_ts_train
## W = 0.9639, p-value = 1.335e-06
```

```
arima <- auto.arima(cpi_ts_train, lambda = lambda)
summary(arima)
```

```
## Series: cpi_ts_train
## ARIMA(2,1,0)(2,0,0)[12] with drift
## Box Cox transformation: lambda= -0.9999242
##
## Coefficients:
##          ar1      ar2      sar1      sar2  drift
##          0.5520 -0.2268  0.1884  0.2045  0e+00
## s.e.      0.0587   0.0581  0.0617  0.0709  1e-04
##
## sigma^2 = 3.709e-09: log likelihood = 2795.75
## AIC=-5579.5   AICc=-5579.2   BIC=-5557.54
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.07664625 1.596256 0.6150133 0.04659053 0.2842446 0.1040759
##              ACF1
## Training set 0.03222512
```

```
checkresiduals(arima)
```

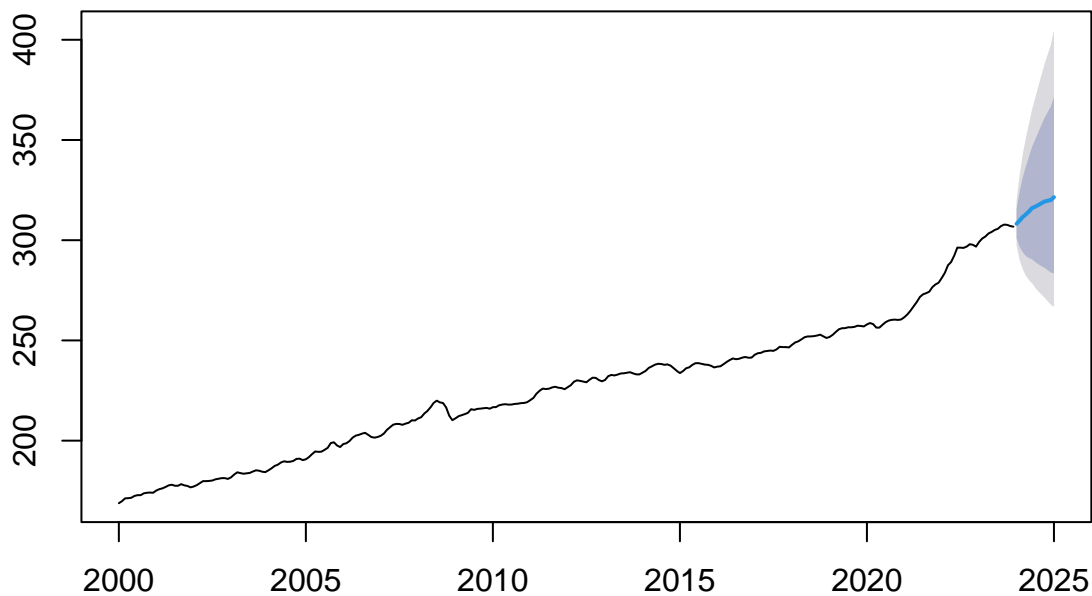


```
##
## Ljung-Box test
```

```
##
## data: Residuals from ARIMA(2,1,0)(2,0,0)[12] with drift
## Q* = 2.0032, df = 20, p-value = 1
##
## Model df: 4. Total lags used: 24
```

```
arima_forecast <- forecast(arima, h = 13)
plot(arima_forecast)
```

### Forecasts from ARIMA(2,1,0)(2,0,0)[12] with drift



```
evaluate_performance(cpi_ts_test, arima_forecast$mean)
```

```
## $MSE
## [1] 7.113968
##
## $RMSE
## [1] 2.667202
##
## $MAE
## [1] 2.194474
##
## $R_squared
## [1] -0.2960262
```

### c. GARCH Model

```
residuals_arima <- residuals(arima)
# Perform ARCH-LM test (using 5 lags)
arch_test <- ArchTest(residuals_arima, lags = 5)
print(arch_test)
```

```
##
## ARCH LM-test; Null hypothesis: no ARCH effects
##
## data: residuals_arima
## Chi-squared = 16.65, df = 5, p-value = 0.005214

# Define GARCH(1,1) Model for CPI
garch_spec <- ugarchspec(variance.model = list(model = "sGARCH", garchOrder = c(1,1)),
  mean.model = list(armaOrder = c(2,2), include.mean = TRUE),
  distribution.model = "norm")

# Fit GARCH model to CPI
garch_model <- ugarchfit(spec = garch_spec, data = cpi_ts_train)

# View results
summary(garch_model)

##      Length      Class      Mode
##          1 uGARCHfit        S4

garch_forecast <- ugarchforecast(garch_model, n.ahead = 13)
evaluate_performance(cpi_ts_test, garch_forecast@forecast$seriesFor)

## $MSE
## [1] 4.751166
##
## $RMSE
## [1] 2.179717
##
## $MAE
## [1] 1.872261
##
## $R_squared
## [1] 0.1344301
```

## (2) Multivariate

```
# Split data into training (2000-01 to 2023-12) and testing (2024-01 onward)

train_end <- which(dataset$Date == as.Date("2023-12-01"))
dataset_train <- dataset[1:train_end, ]
dataset_test <- dataset[(train_end + 1):nrow(dataset), ]

dataset_lag_train <- dataset_lag[1:train_end, ]
dataset_lag_test <- dataset_lag[(train_end + 1):nrow(dataset), ]

cor(dataset_train$CPI, dataset_train$FEDFUNDS, use = "complete.obs")

## [1] -0.2372331
```



```
cor(dataset_train$CPI, dataset_train$UNRATE, use = "complete.obs")
```

```
## [1] -0.1512589
```

```
cor(dataset_train$CPI, dataset_train$M2, use = "complete.obs")
```

```
## [1] 0.9661038
```

```
granger_test <- grangertest(CPI ~ FEDFUNDS, order = 2, data = dataset_train)
print("FEDFUNDS:")
```

```
## [1] "FEDFUNDS:"
```

```
print(granger_test)
```

```
## Granger causality test
##
## Model 1: CPI ~ Lags(CPI, 1:2) + Lags(FEDFUNDS, 1:2)
## Model 2: CPI ~ Lags(CPI, 1:2)
##   Res.Df Df       F Pr(>F)
## 1     281
## 2     283 -2 0.9272 0.3969
```

```
granger_test <- grangertest(CPI ~ UNRATE, order = 2, data = dataset_train)
print("UNRATE:")
```

```
## [1] "UNRATE:"
```

```
print(granger_test)
```

```
## Granger causality test
##
## Model 1: CPI ~ Lags(CPI, 1:2) + Lags(UNRATE, 1:2)
## Model 2: CPI ~ Lags(CPI, 1:2)
##   Res.Df Df       F Pr(>F)
## 1     281
## 2     283 -2 0.0448 0.9561
```

```
granger_test <- grangertest(CPI ~ M2, order = 2, data = dataset_train)
print("M2:")
```

```
## [1] "M2:"
```

```
print(granger_test)
```

```
## Granger causality test
##
## Model 1: CPI ~ Lags(CPI, 1:2) + Lags(M2, 1:2)
## Model 2: CPI ~ Lags(CPI, 1:2)
##   Res.Df Df       F    Pr(>F)
## 1      281
## 2      283 -2 8.3245 0.0003074 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Reg-ARMA

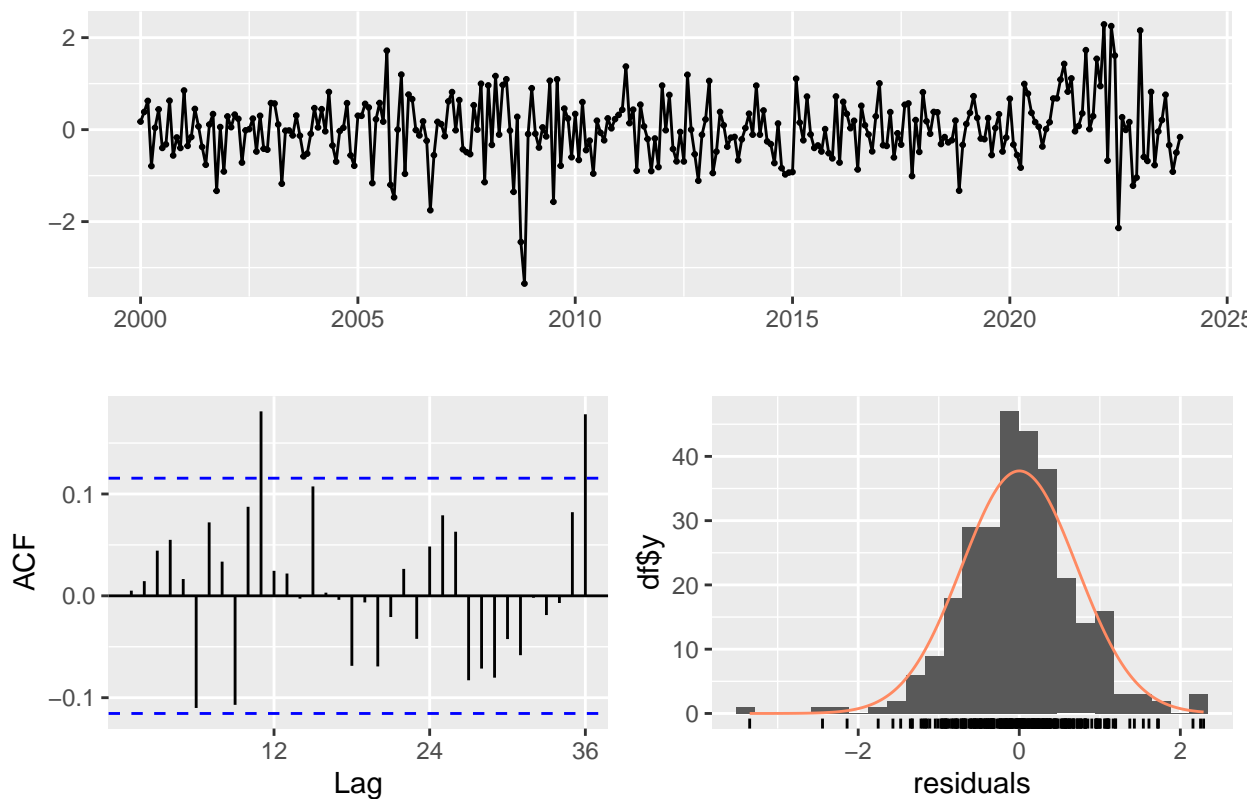
```
# Regression with ARMA Errors
reg_arma <- auto.arima(cpi_ts_train,
                      xreg = as.matrix(dataset_train[, c("M2")]),
                      seasonal = TRUE)

summary(reg_arma)
```

```
## Series: cpi_ts_train
## Regression with ARIMA(0,1,2)(0,0,2)[12] errors
##
## Coefficients:
##          ma1      ma2      sma1      sma2      drift      xreg
##          0.6040  0.1331  0.1931  0.1378  0.5519 -0.0013
## s.e.      0.0592  0.0621  0.0628  0.0578  0.1014  0.0006
##
## sigma^2 = 0.5222:  log likelihood = -311.54
## AIC=637.08   AICc=637.48   BIC=662.69
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.0007647766 0.7138165 0.5302602 -0.004068395 0.2349103 0.08973349
##              ACF1
## Training set 0.005092052
```

```
# Check residuals of the model
checkresiduals(reg_arma)
```

## Residuals from Regression with ARIMA(0,1,2)(0,0,2)[12] errors



```
##
##  Ljung-Box test
##
## data:  Residuals from Regression with ARIMA(0,1,2)(0,0,2)[12] errors
## Q* = 31.194, df = 20, p-value = 0.05268
##
## Model df: 4.   Total lags used: 24
```

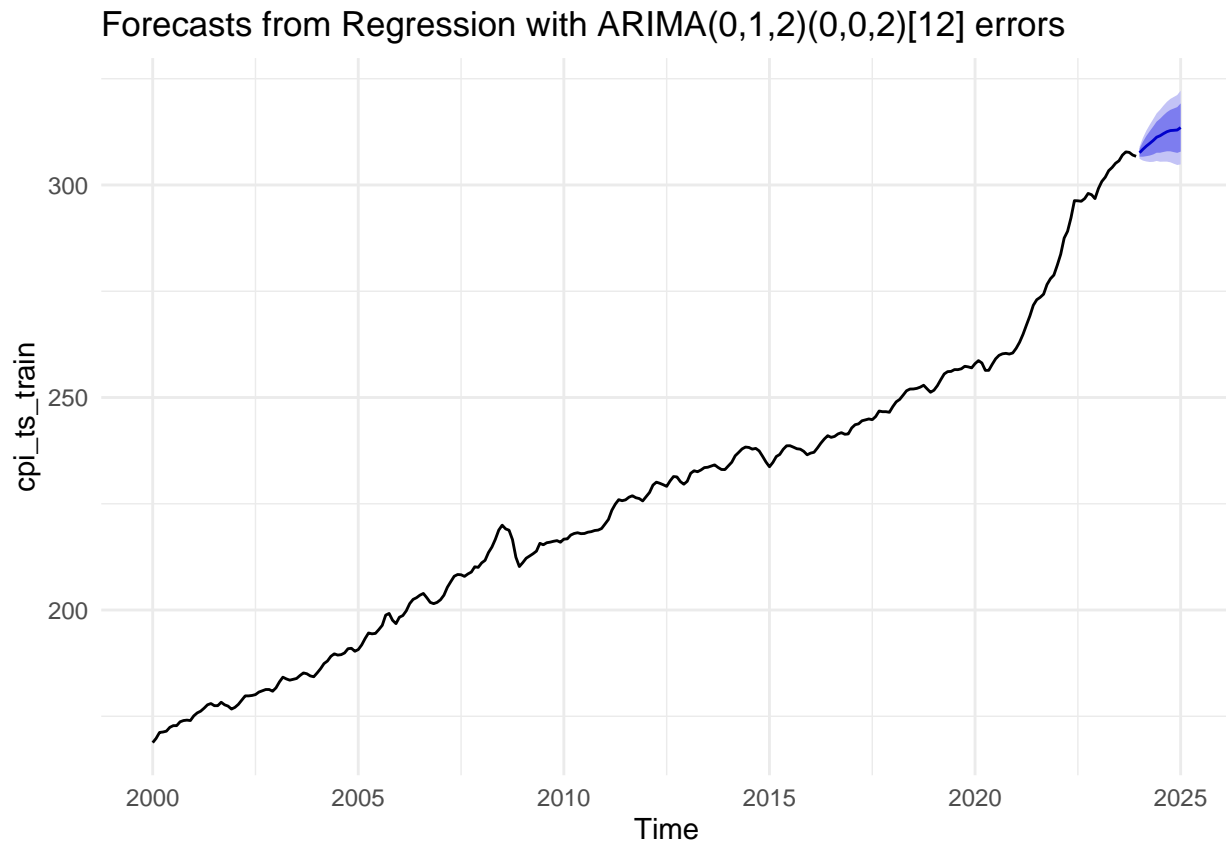
```
# Make predictions
arma_m2 <- auto.arima(dataset_train[, c("M2")], lambda = "auto")
m2_forecast <- forecast(arma_m2, h = 13)
```

```
reg_arma_forecast <- forecast(reg_arma,
                              xreg = as.matrix(m2_forecast$mean),
                              h = nrow(dataset_lag_test))
evaluate_performance(cpi_ts_test, reg_arma_forecast$mean)
```

```
## $MSE
## [1] 8.691107
##
## $RMSE
## [1] 2.948068
##
## $MAE
## [1] 2.838863
```

```
##
## $R_squared
## [1] -0.5833501
```

```
autoplot(reg_arma_forecast)+theme_minimal()
```

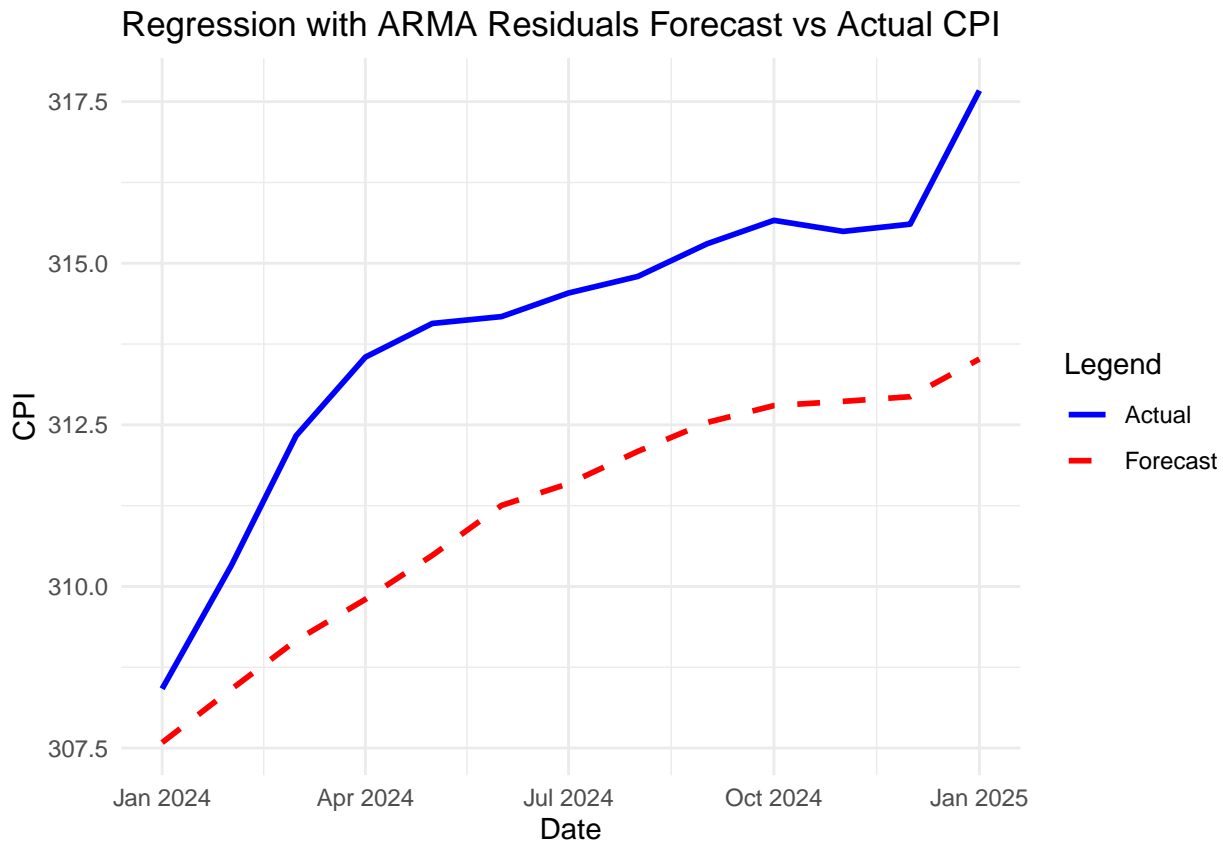


```
# Create a dataframe for plotting
plot_data <- data.frame(
  Date = dataset_test$Date,
  Actual = cpi_ts_test,
  Predicted = reg_arma_forecast$mean
)

# Plot forecast vs actual
ggplot(plot_data, aes(x = Date)) +
  geom_line(aes(y = Actual, color = "Actual"), size = 1) +
  geom_line(aes(y = Predicted, color = "Forecast"), size = 1, linetype = "dashed") +
  labs(title = "Regression with ARMA Residuals Forecast vs Actual CPI", x = "Date", y = "CPI") +
  scale_color_manual(name = "Legend", values = c("Actual" = "blue", "Forecast" = "red")) +
  theme_minimal()
```

```
## Warning: Using 'size' aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use 'linewidth' instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
## to continuous.
```



## VAR

```
# Check stationarity of each variable
adf.test(diff(log(dataset_train$CPI)))
```

```
## Warning in adf.test(diff(log(dataset_train$CPI))): p-value smaller than printed
## p-value
```

```
##
## Augmented Dickey-Fuller Test
##
## data: diff(log(dataset_train$CPI))
## Dickey-Fuller = -6.4423, Lag order = 6, p-value = 0.01
## alternative hypothesis: stationary
```

```
adf.test(diff(log(dataset_train$M2)))
```

```
## Warning in adf.test(diff(log(dataset_train$M2))): p-value smaller than printed
## p-value
```

```

##
## Augmented Dickey-Fuller Test
##
## data: diff(log(dataset_train$M2))
## Dickey-Fuller = -4.3114, Lag order = 6, p-value = 0.01
## alternative hypothesis: stationary

# Log transformation and differencing
dataset_train$CPI_log_diff <- c(NA, diff(log(dataset_train$CPI)))
dataset_test$CPI_log_diff <- c(NA, diff(log(dataset_test$CPI)))

# Log transformation and differencing on M2
dataset_train$M2_log_diff <- c(NA, diff(log(dataset_train$M2)))
dataset_test$M2_log_diff <- c(NA, diff(log(dataset_test$M2)))

var_train <- cbind(
  dataset_train$CPI_log_diff,
  dataset_train$M2_log_diff
)

var_train <- na.omit(var_train)
colnames(var_train) <- c("CPI_log_diff", "M2_log_diff")

var_model <- VAR(var_train, p = 10, type = "both", season = 12)
var_forecast <- predict(var_model, n.ahead = nrow(dataset_test))

# Extract CPI forecasts and convert back from log differences
cpi_forecast_log_diff <- var_forecast$fcst$CPI_log_diff[,1] # Forecasted log differences
cpi_forecast <- exp(log(tail(dataset_train$CPI, 1)) + cumsum(cpi_forecast_log_diff))
print(evaluate_performance(dataset_test$CPI, cpi_forecast))

## $MSE
## [1] 2.068678
##
## $RMSE
## [1] 1.43829
##
## $MAE
## [1] 1.262145
##
## $R_squared
## [1] 0.6231271

# Create dataframe for plotting
plot_data <- data.frame(
  Date = time(cpi_ts_test),
  Actual = dataset_test$CPI,
  Forecast = cpi_forecast
)

# Plot actual vs forecasted CPI
ggplot(plot_data, aes(x = Date)) +
  geom_line(aes(y = Actual, color = "Actual"), size = 1) +

```

```
geom_line(aes(y = Forecast, color = "Forecast"), size = 1, linetype = "dashed") +
labs(title = "VAR Model Forecast vs Actual CPI", x = "Time", y = "CPI") +
scale_color_manual(name = "Legend", values = c("Actual" = "blue", "Forecast" = "red")) +
theme_minimal()
```

## Don't know how to automatically pick scale for object of type <ts>. Defaulting  
## to continuous.

